

# Introduction to SQL on Apache Flink<sup>®</sup>

---

Flink SQL Training

<https://github.com/ververica/sql-training>

# Motivation

# Flink's Powerful Abstractions

Layered abstractions to  
navigate simple to complex use cases

High-level  
Analytics API

**SQL / Table API (dynamic tables)**

```
SELECT room, TUMBLE_END(rowtime, INTERVAL '1' HOUR), AVG(temp)
FROM sensors
GROUP BY TUMBLE(rowtime, INTERVAL '1' HOUR), room
```



Stream- & Batch  
Data Processing

**DataStream API (streams, windows)**

```
val stats = stream
  .keyBy("sensor")
  .timeWindow(Time.seconds(5))
  .sum((a, b) -> a.add(b))
```



Stateful Event-  
Driven Applications

**Process Function (events, state, time)**

```
def processElement(event: MyEvent, ctx: Context, out: Collector[Result]) = {
  // work with event and state
  (event, state.value) match { ... }

  out.collect(...) // emit events
  state.update(...) // modify state

  // schedule a timer callback
  ctx.timerService.registerEventTimeTimer(event.timestamp + 500)
}
```



# The DataStream API is great...

- Very expressive stream processing API
  - Transform, aggregate, and join events
  - Java and Scala
- Control how events are processed with respect to time
  - Timestamps, Watermarks, Windows, Timers, Triggers, Allowed Lateness, ...
- Maintain and update application state
  - Keyed state, operator state, state backends, checkpointing, ...



# ... but it's not made for everyone.

- Writing distributed programs is not always easy
  - Stream processing technology spreads rapidly
  - New concepts (time, state, ...)
- Requires knowledge & skill
  - Continuous applications have special requirements
  - Programming experience (Java / Scala)
- Users want to focus on their business logic



# Why not SQL (or another relational API)?

- Relational APIs are declarative
  - User says what is needed, system decides how to compute it
- Queries can be effectively optimized
  - Less imperative black-box code
  - Well-researched field
- Queries are efficiently executed
  - Let Flink deal with state and time
- “Everybody” knows and uses SQL



# Goals

- Easy, declarative, and concise relational API
- Expressive enough for a wide range of use cases
- Unified syntax and semantics for batch & streaming data



# Table API & SQL



# Apache Flink's Relational APIs

## ANSI SQL

```
SELECT user, COUNT(url) AS cnt  
FROM clicks  
GROUP BY user
```

## LINQ-style Table API

```
tableEnvironment  
    .scan("clicks")  
    .groupBy('user')  
    .select('user', 'url.count as 'cnt')
```

**Unified APIs for batch & streaming data**

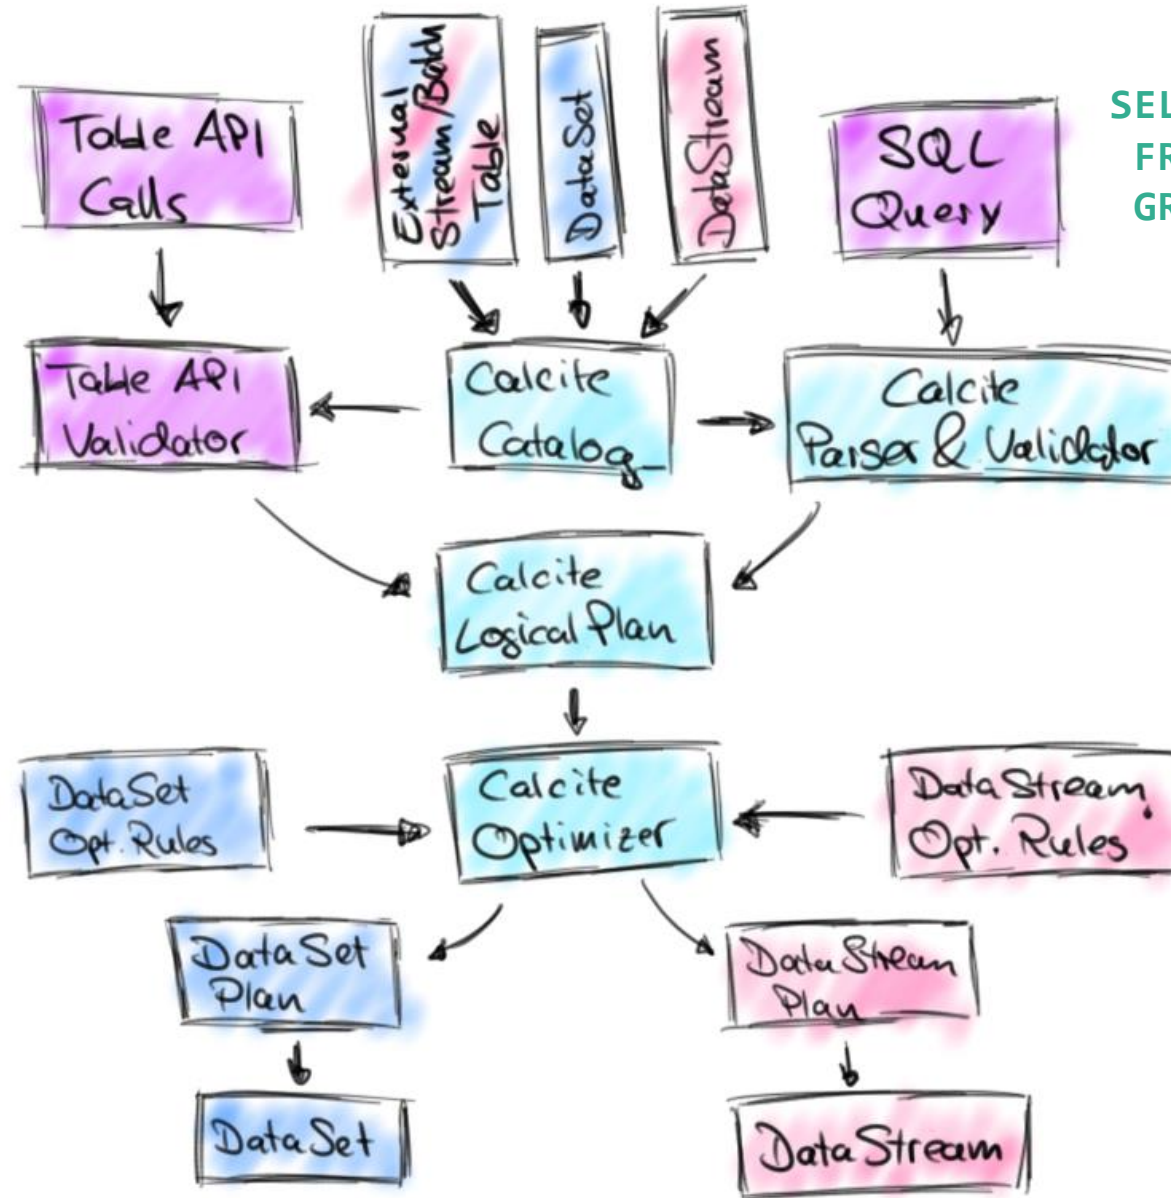
***A query specifies exactly the same result  
regardless whether its input is  
static batch data or streaming data.***



# Query Translation

```
tableEnvironment  
  .scan("clicks")  
  .groupBy('user')  
  .select('user', 'url.count')
```

```
SELECT user, COUNT(url) AS cnt  
FROM clicks  
GROUP BY user
```

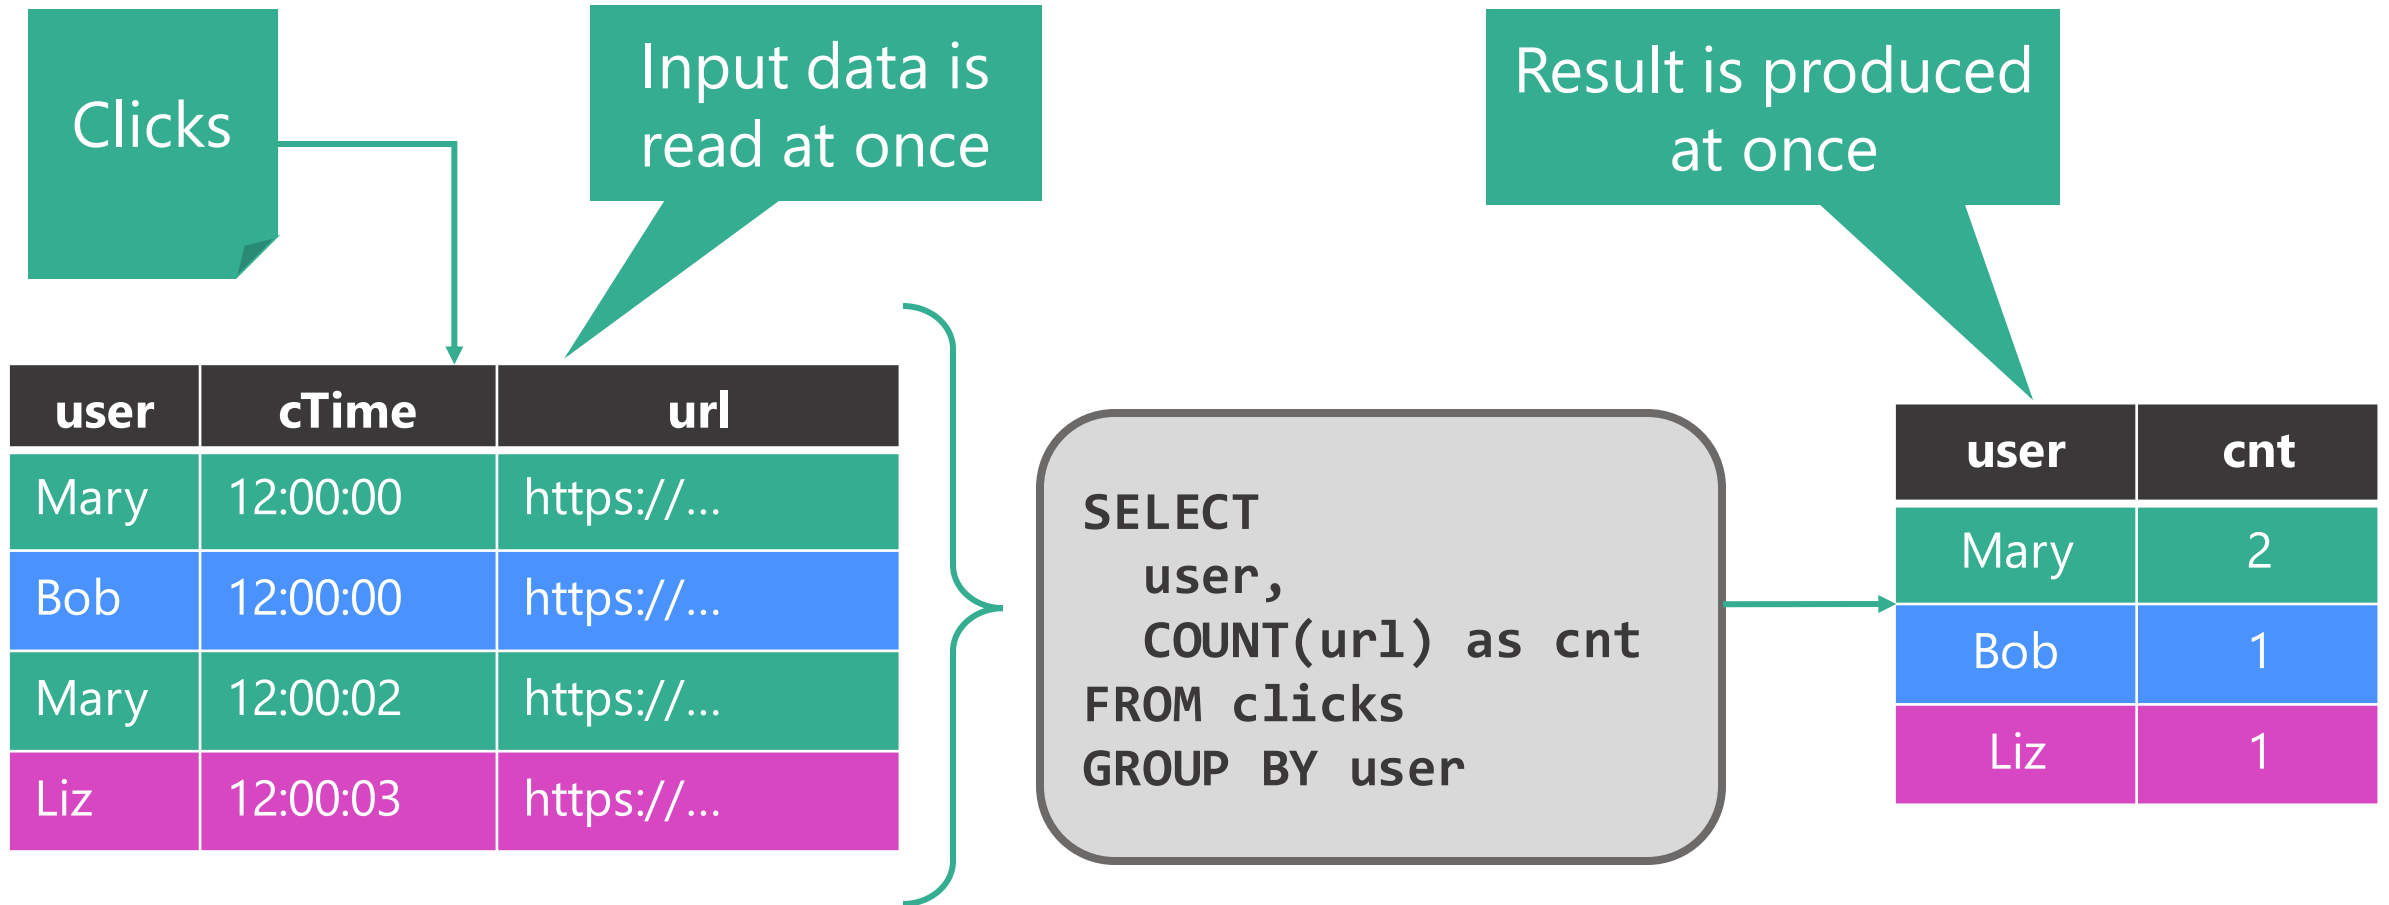


Input data is  
bounded  
(batch)

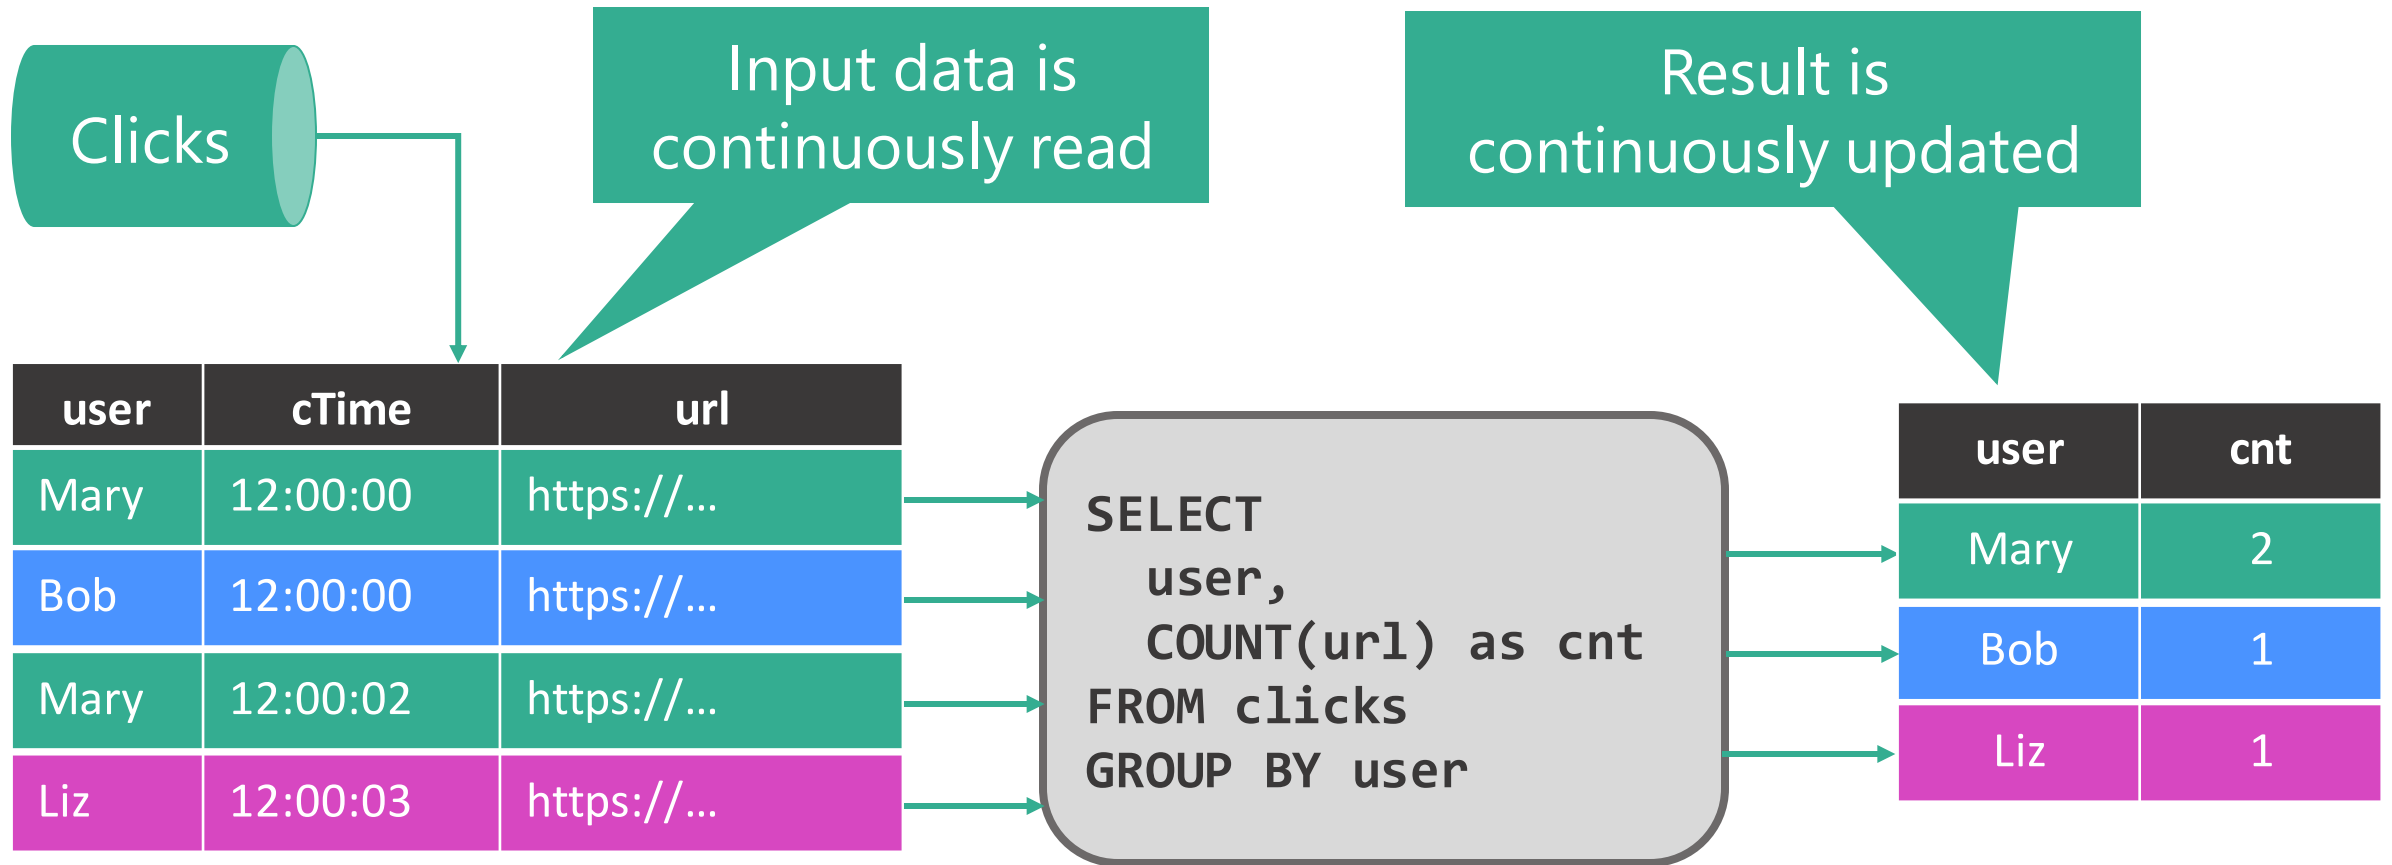
Input data is  
unbounded  
(streaming)



# What if “Clicks” is a File?



# What if “Clicks” is a Stream?



**The result is the same!**

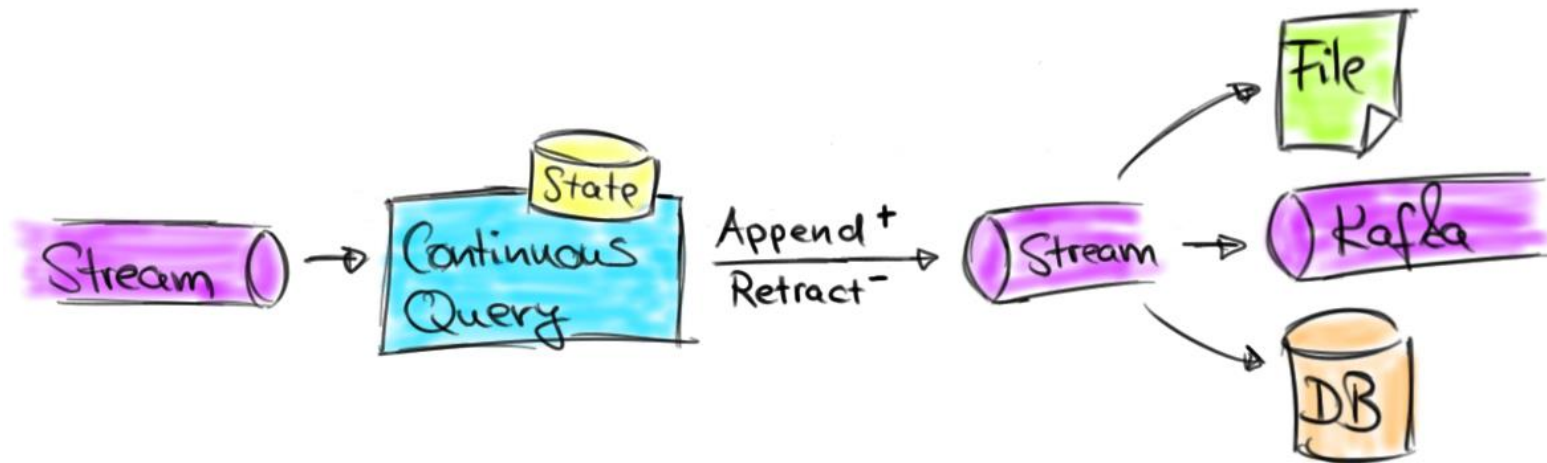
# Use Cases

# Flink SQL is Used in Production



# Data Pipelines

- Transform, aggregate, and move events in real-time
- Low-latency ETL
  - Convert and write streams to file systems, DBMS, K-V stores, indexes, ...
  - Ingest appearing files to produce streams



# Data Pipelines

- Support for POJOs, maps, arrays, and other nested types
- Large set of built-in functions (150+)
  - LIKE, EXTRACT, TIMESTAMPADD, FROM\_BASE64, MD5, STDDEV\_POP, AVG, ...
- Support for custom UDFs (scalar, table, aggregate)

See also:

<https://ci.apache.org/projects/flink/flink-docs-stable/dev/table/functions.html>


<https://ci.apache.org/projects/flink/flink-docs-stable/dev/table/udfs.html>





# Stream & Batch Analytics

- Stream & Batch Analytics
  - Run analytical queries over bounded and unbounded data
  - Query and compare historic and real-time data
  - Compute and update data to visualize in real-time



# SQL Feature Set in Flink 1.9

## **STREAMING & BATCH**

- SELECT FROM WHERE
- GROUP BY [HAVING]
  - Non-windowed
  - TUMBLE, HOP, SESSION windows
- JOIN
  - Time-Windowed INNER + OUTER JOIN
  - Non-windowed INNER + OUTER JOIN
- User-Defined Functions
  - Scalar
  - Aggregation
  - Table-valued

## **STREAMING ONLY**

- OVER / WINDOW
  - UNBOUNDED / BOUNDED PRECEDING
- INNER JOIN with time-versioned table
- MATCH\_RECOGNIZE
  - Pattern Matching/CEP (SQL:2016)

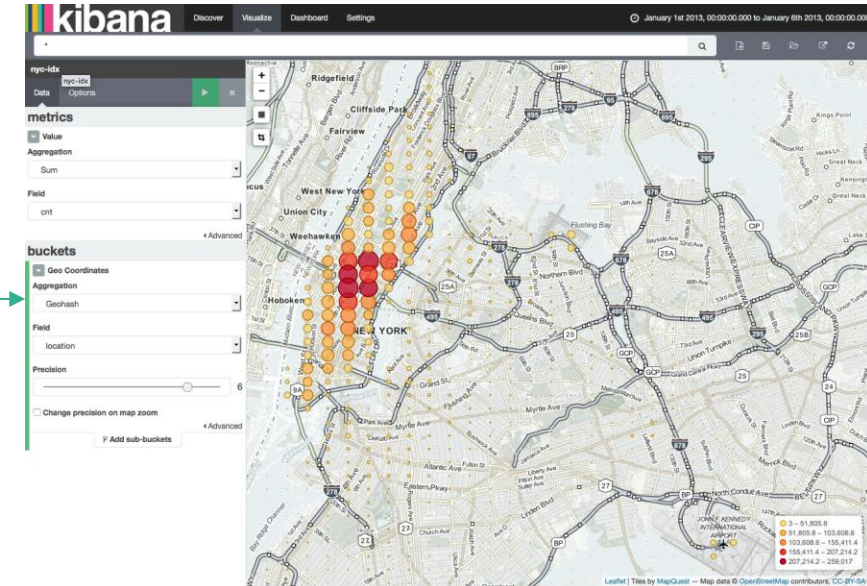
## **BATCH ONLY**

- UNION / INTERSECT / EXCEPT
- ORDER BY



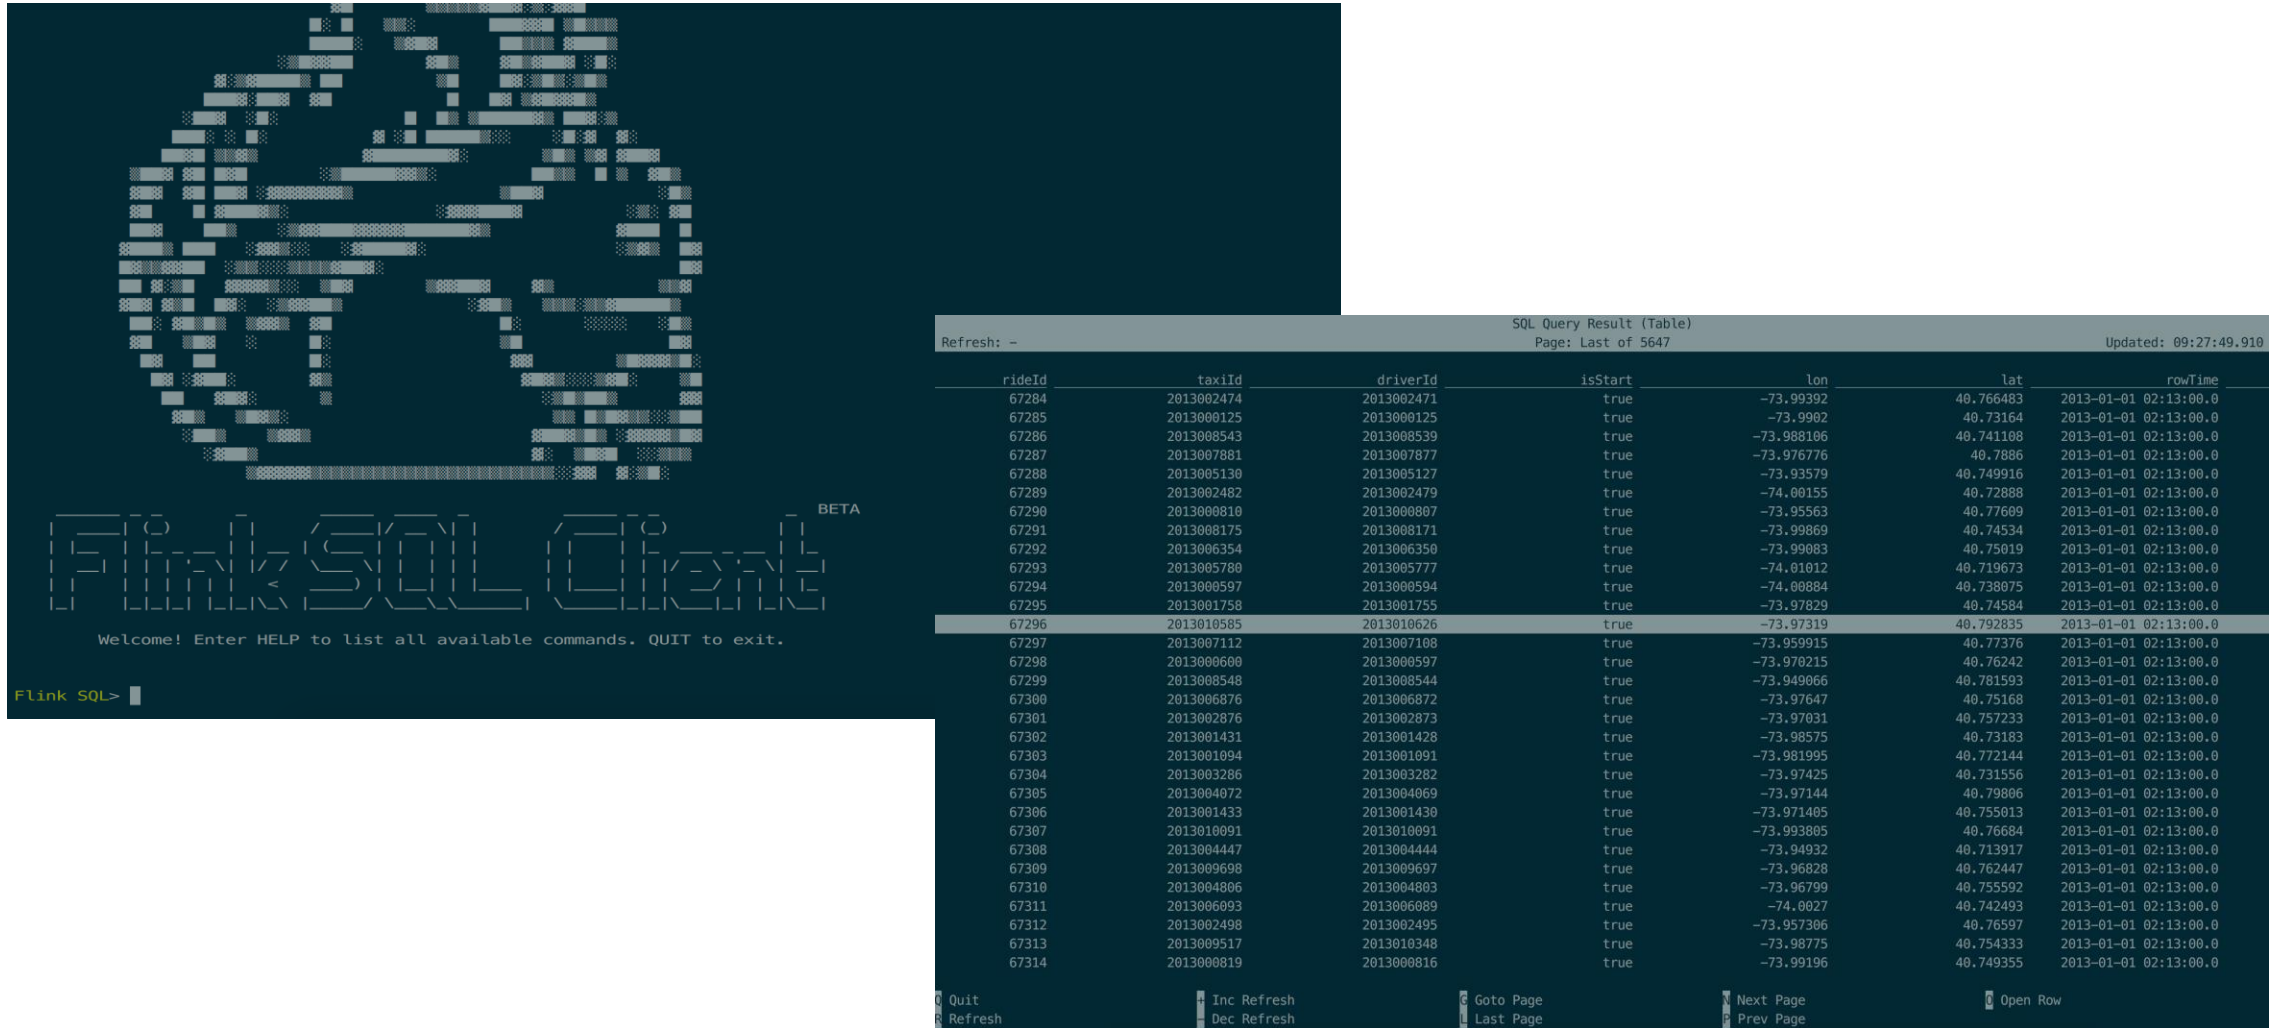
# Building a Dashboard

```
SELECT cell,  
       isStart,  
       HOP_END(rowtime, INTERVAL '5' MINUTE, INTERVAL '15' MINUTE) AS hopEnd,  
       COUNT(*) AS cnt  
FROM (SELECT rowtime, isStart, toCellId(lon, lat) AS cell  
      FROM TaxiRides)  
GROUP BY cell,  
       isStart,  
       HOP(rowtime, INTERVAL '5' MINUTE, INTERVAL '15' MINUTE)
```



SQL Client

# Introduction to SQL Client



The image displays the Flink SQL Client interface. On the left, a terminal window shows a large ASCII art logo of a globe with the text "Flink SQL Client" and "BETA" below it. The terminal prompt is "Flink SQL>". Below the logo, it says "Welcome! Enter HELP to list all available commands. QUIT to exit."

On the right, a table displays the results of an SQL query. The table has 8 columns: rideId, taxiId, driverId, isStart, lon, lat, and rowTime. The table is titled "SQL Query Result (Table)" and "Page: Last of 5647". The table is updated at "09:27:49.910".

rideId	taxiId	driverId	isStart	lon	lat	rowTime
67284	2013002474	2013002471	true	-73.99392	40.766483	2013-01-01 02:13:00.0
67285	2013000125	2013000125	true	-73.9902	40.73164	2013-01-01 02:13:00.0
67286	2013000543	2013000539	true	-73.988106	40.741108	2013-01-01 02:13:00.0
67287	2013007881	2013007877	true	-73.976776	40.7886	2013-01-01 02:13:00.0
67288	2013005130	2013005127	true	-73.93579	40.749916	2013-01-01 02:13:00.0
67289	2013002482	2013002479	true	-74.00155	40.72888	2013-01-01 02:13:00.0
67290	2013000810	2013000807	true	-73.95563	40.77609	2013-01-01 02:13:00.0
67291	2013008175	2013008171	true	-73.99869	40.74534	2013-01-01 02:13:00.0
67292	2013006354	2013006350	true	-73.99083	40.75019	2013-01-01 02:13:00.0
67293	2013005780	2013005777	true	-74.01012	40.719673	2013-01-01 02:13:00.0
67294	2013000597	2013000594	true	-74.00884	40.738075	2013-01-01 02:13:00.0
67295	2013001758	2013001755	true	-73.97829	40.74584	2013-01-01 02:13:00.0
67296	2013010585	2013010626	true	-73.97319	40.792835	2013-01-01 02:13:00.0
67297	2013007112	2013007108	true	-73.959015	40.77376	2013-01-01 02:13:00.0
67298	2013000600	2013000597	true	-73.970215	40.76242	2013-01-01 02:13:00.0
67299	2013008548	2013008544	true	-73.949066	40.781593	2013-01-01 02:13:00.0
67300	2013006876	2013006872	true	-73.97647	40.75168	2013-01-01 02:13:00.0
67301	2013002876	2013002873	true	-73.97031	40.757233	2013-01-01 02:13:00.0
67302	2013001431	2013001428	true	-73.98575	40.73183	2013-01-01 02:13:00.0
67303	2013001094	2013001091	true	-73.981995	40.772144	2013-01-01 02:13:00.0
67304	2013003286	2013003282	true	-73.97425	40.731556	2013-01-01 02:13:00.0
67305	2013004072	2013004069	true	-73.97144	40.79806	2013-01-01 02:13:00.0
67306	2013001433	2013001430	true	-73.971405	40.755013	2013-01-01 02:13:00.0
67307	2013010091	2013010091	true	-73.993805	40.76684	2013-01-01 02:13:00.0
67308	2013004447	2013004444	true	-73.94032	40.713917	2013-01-01 02:13:00.0
67309	2013009698	2013009697	true	-73.96828	40.762447	2013-01-01 02:13:00.0
67310	2013004806	2013004803	true	-73.96799	40.755592	2013-01-01 02:13:00.0
67311	2013006093	2013006089	true	-74.0027	40.742493	2013-01-01 02:13:00.0
67312	2013002498	2013002495	true	-73.957306	40.76597	2013-01-01 02:13:00.0
67313	2013009517	2013010348	true	-73.98775	40.754333	2013-01-01 02:13:00.0
67314	2013000819	2013000816	true	-73.99196	40.749355	2013-01-01 02:13:00.0

At the bottom of the table, there are navigation controls: Quit, Refresh, Inc Refresh, Dec Refresh, Goto Page, Last Page, Next Page, Prev Page, and Open Row.



# Introduction to SQL Client

- Flink without a single line of code
  - Only SQL and YAML
  - Add connectors and formats by downloading SQL JAR files
- Use cases
  - Query prototyping
  - Ad-hoc stream analytics & inspection
  - Detached query submission



# SQL Client Environment Files

- Non-programmatic way of configuring Flink jobs
- Per-session and/or global configuration in YAML
- Environment file defines
  - Table schema and connection details to external systems
  - Views
  - User-defined functions
  - Execution properties (e.g. result mode, execution mode)
  - Deployment properties



# SQL Client Environment File Example

```
1 # Define table sources and sinks here.
2 tables:
3   -- name: MyTableSource
4   -- type: source
5   -- update-mode: append
6   -- connector:
7     -- type: filesystem
8     -- path: "/path/to/something.csv"
9   -- format:
10    -- type: csv
11    -- fields:
12      -- name: MyField1
13      -- type: INT
14      -- name: MyField2
15      -- type: VARCHAR
16    -- line-delimiter: "\n"
17    -- comment-prefix: "#"
18  -- schema:
19    -- name: MyField1
20    -- type: INT
21    -- name: MyField2
22    -- type: VARCHAR
23
24 # Define table views here.
25 views:
26   -- name: MyCustomView
27   -- query: "SELECT MyField2 FROM MyTableSource"
28
29 # Define user-defined functions here.
30 functions:
31   -- name: myUDF
32   -- from: class
33   -- class: foo.bar.AggregateUDF
34
35 # Execution properties allow for changing the behavior of a table program.
36 execution:
37   -- type: streaming .....# required: execution mode either 'batch' or 'streaming'
38   -- result-mode: table .....# required: either 'table' or 'changelog'
39   -- parallelism: 1 .....# optional: Flink's parallelism (1 by default)
```

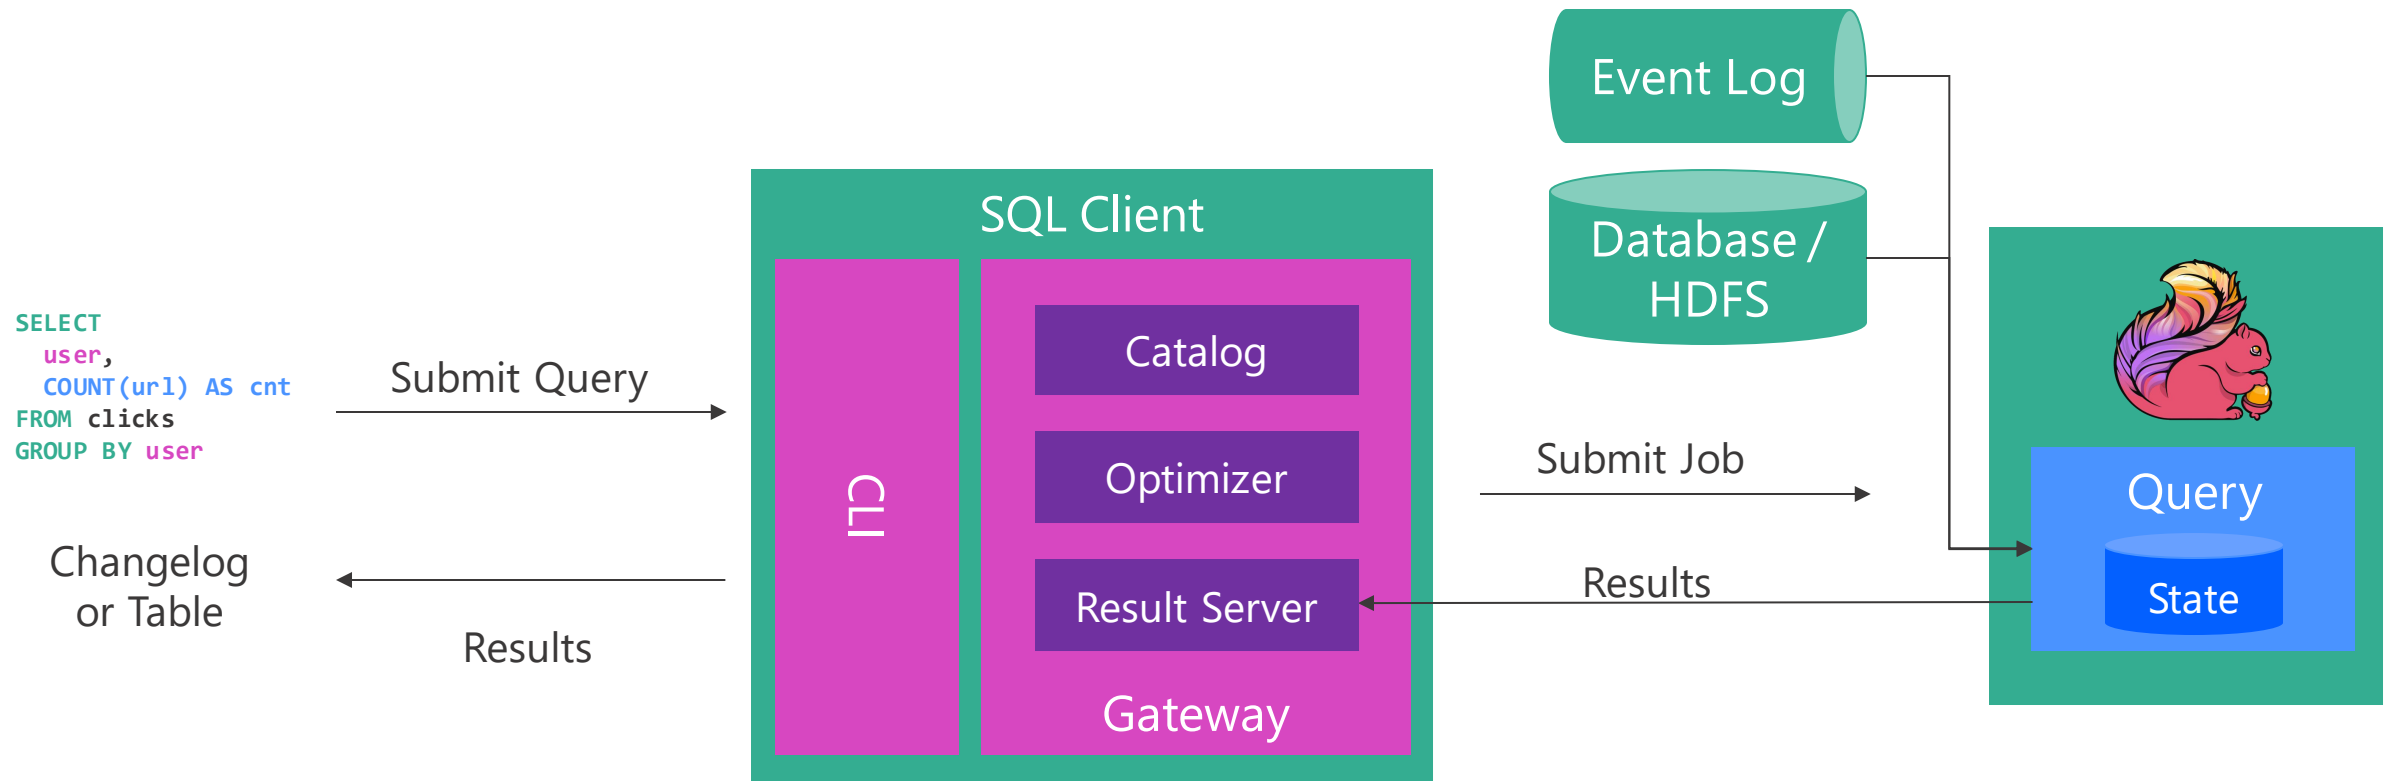
See also:

<https://ci.apache.org/projects/flink/flink-docs-stable/dev/table/sqlClient.html>

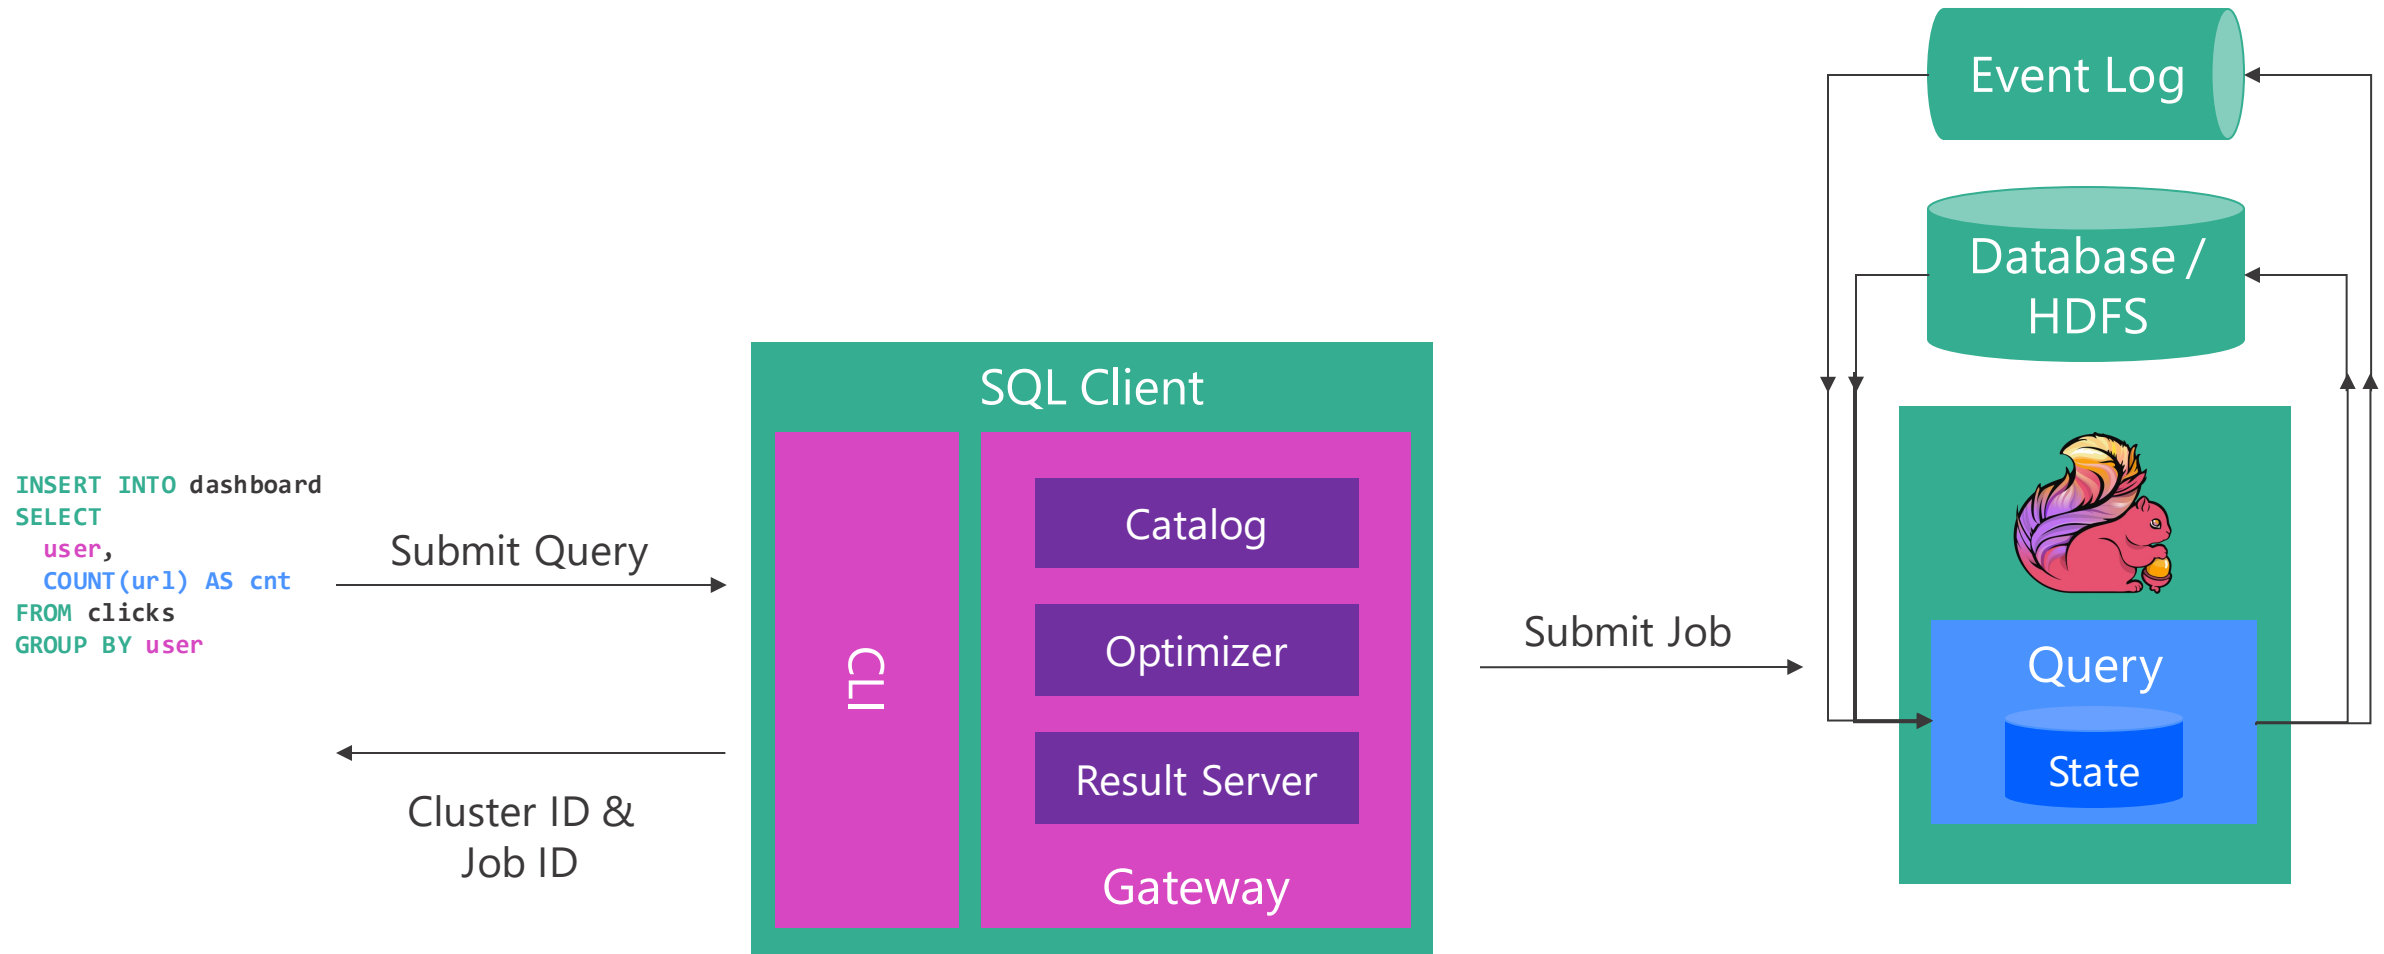




# Interactive Query Submission via SQL Client



# Detached Query Submission via SQL Client



# Hands On Exercises

# Introduction to SQL on Flink

Continue with the “Introduction to the Training Environment”  
in “Introduction to SQL on Flink”

<https://github.com/ververica/sql-training/wiki/Introduction-to-SQL-on-Flink>

We are here to help!





# ververica

---

[www.ververica.com](http://www.ververica.com)

@VervericaData